

Distributed Log Aggregator - Milestone 2 Documentation

Prisacaru Bogdan-Paul, 343C4

December 14, 2025

1 Overview

The **Distributed Log Aggregator** is a scalable, microservices-based log aggregation and analytics platform. It is built using **Docker Swarm**, **FastAPI**, **Keycloak**, and **OpenSearch**.

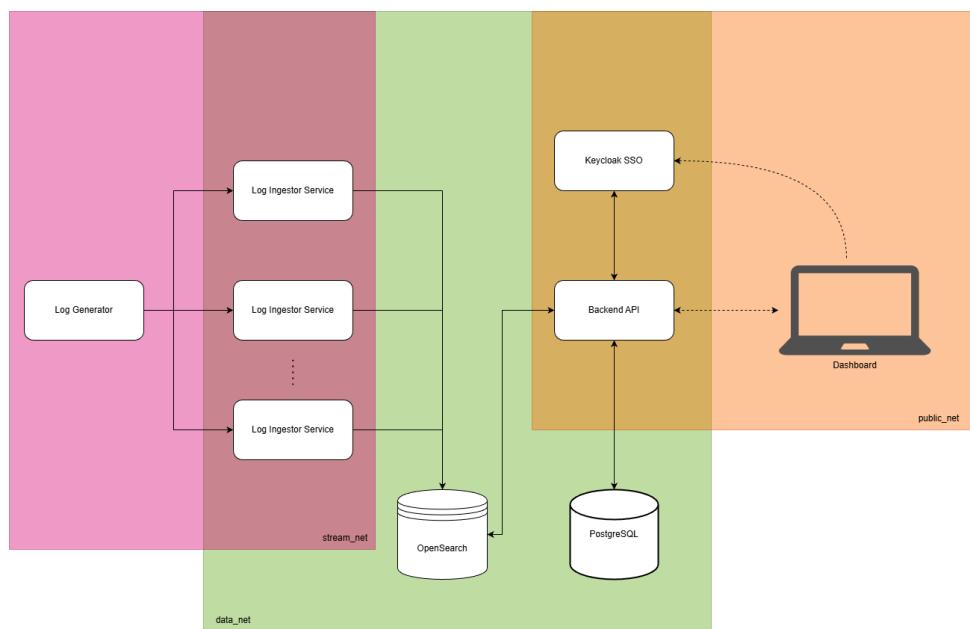


Figure 1: Milestone 2 System Architecture Overview

2 Architecture

The system is composed of the following microservices:

- **Dashboard Service**: A FastAPI-based frontend and API providing a user interface for log visualization, filtering, and exporting. It handles authentication via Keycloak.
- **Ingestor Service**: A lightweight service dedicated to receiving logs from various sources and indexing them into OpenSearch.
- **Log Generator**: A utility service that generates mock log traffic to simulate a real-world distributed system.
- **OpenSearch**: The core search and analytics engine used to store and query logs.
- **Keycloak**: Handles Identity and Access Management (IAM) and Single Sign-On (SSO).
- **PostgreSQL**: Stores user profile data and application-specific metadata.

3 Key Features

3.1 Authentication

- Single Sign-On (SSO): Powered by Keycloak (OpenID Connect).

3.2 Role-Based Access Control (RBAC)

The system implements secure access management with distinct roles. The permissions are defined as follows:

Feature	Admin	Developer	Viewer
Login via SSO	✓	✓	✓
Search & Filter	✓	✓	✓
Time Window	Unlimited	Unlimited	Last 3 Hours
Log Visibility	All Levels	All Levels	INFO/WARN Only
Export Data	✓	✓	✗

3.3 Advanced Filtering & Search

- Full-Text Search: Search log messages using OpenSearch.
- Structured Filters: Filter by **Service**, **Log Level**, and **Time Range**.
- Pagination: Efficiently browse through logs.

3.4 Data Export

- Export filtered logs to **CSV** or **JSON** formats.
- The export functionality respects active filters (e.g., filtering for "ERROR" results in an export containing only errors).

3.5 Scalable Ingestion

- Implements a time-based indexing strategy (`app-logs-YYYY.MM.DD`) for efficient storage management.

4 Setup & Deployment

4.1 Prerequisites

- Docker & Docker Compose
- Docker Swarm initialized (`docker swarm init`)

4.2 Deployment Steps

1. Build the Services:

```
docker compose build
```

2. Deploy the Stack:

```
docker stack deploy -c docker-compose.yml log_stack
```

3. Access the Dashboard: Open <http://localhost:8000> in your browser.

4. Access Keycloak Console: Open <http://localhost:8080>.

5 Default Credentials

5.1 Application Users (Log Realm)

Role	Username	Password
Admin	admin	admin
Developer	developer	developer
Viewer	viewer	viewer

5.2 Keycloak Administration Console

- URL: <http://localhost:8080>
- Username: admin
- Password: admin

6 Testing

To generate traffic, ensure the log generator is running:

```
docker service scale log_stack_log-generator=1
```

To stop generation:

```
docker service scale log_stack_log-generator=0
```

6.1 Automated Testing with Postman

A comprehensive Postman collection is provided in `tests/postman_collection.json` to validate the system's functionality and security requirements.

The collection includes tests for:

- **Authentication:** Verifying login flows for Admin, Developer, and Viewer roles using Keycloak's API.
- **RBAC Enforcement:** Ensuring that restricted endpoints (e.g., `/export`) are accessible only to authorized roles.
- **Log Ingestion:** Verifying that logs are correctly accepted by the Ingestor Service.
- **Search Functionality:** Validating that ingested logs are indexed and searchable via the Dashboard.

To run the tests:

1. Import the `postman_collection.json` file into Postman.
2. Ensure the stack is running locally.
3. Run the collection using the Postman Collection Runner.

7 Project Repository

The complete source code for this milestone is available on GitHub:

<https://github.com/BogdanPaul15/Distributed-Log-Aggregator>